

AS per the Requirement Document:

- Blocker

Why it's blocked

1. Qwen3-14B/32B, Jais models :
.

our laptop's RAM/VRAM (~7.6GB) is far below what these need (16-96GB)

2. Reliable phone/mobile automation:
.

WSL's USB passthrough is unreliable for ADB/Appium

3. Live YouTube/Brave discovery:
.hardware

Needs real API keys/accounts, regardless of

4. MinIO's/S3:

licensing review before choosing it for proprietary production deployment.

5. some commercial fingerprint databases/services are paid

CRM APIS : allows developers to programmatically read, write, and sync customer data between a **Customer Relationship Management platform** and external systems. It connects your central customer database to mobile apps, websites, accounting tools, or databases to automate repetitive workflows.

Qwen3 exact checkpoint, Jais exact checkpoint, MinIO edition/use, YouTube API, Brave Search API, Fingerprinting technology, Docker usage/licensing:

17.1 Hardware procurement

Item	Quantity	Minimum requirement	Recommended
GPU inference server	1	24–32 CPU cores, 128 GB ECC, 32 GB VRAM, 8 TB combined NVMe/evidence capacity	256 GB ECC, 48–96 GB VRAM, mirrored enterprise NVMe, 24 TB+ object storage
Android device gateway	1	8 cores, 32 GB RAM, 1 TB SSD	Business workstation/mini-PC with wired LAN and UPS
Powered USB hub	1	USB 3, adequate total power	Individually switchable ports and surge protection
Android phones	2 initially	Android 11+, 128 GB, current patches	Two manufacturers/device classes; expand only after stable workflow
iOS Mac gateway	0 or 1	Current supported macOS/Xcode, 16 GB RAM	24–32 GB RAM, 1 TB SSD

Confidential technical [handoff](#) | 38

SBA/FIT AI Infrastructure and Mobile Scanning Runbook

Item	Quantity	Minimum requirement	Recommended
iPhones	As approved	Supported iOS and 128 GB	Dedicated non-personal devices
Storage/backup	At least 2 targets	8–16 TB active plus off-site copy	24–40 TB active, versioned, separate off-site replication
Networking	1	1 GbE and private VPN	10 GbE server/storage, dedicated device VLAN
UPS	2	Server and gateway protection	Network-managed UPS with graceful shutdown

2 Software/license review


- Approve the exact LLM checkpoint license for government/commercial use; do not approve only the model family name.
- Record whether a model is Apache 2.0, custom open-weight, non-commercial, or subject to an acceptable-use policy.
- Review App Store/Play Store terms and account use for automated evidence collection.
- Confirm YouTube Data API and Brave Search API account/quota terms.
- Approve Apple Developer membership/signing for iOS automation.
- Confirm MinIO/S3, database, queue, monitoring, and VPN products meet organizational support requirements.
- Document support ownership for consumer versus data-center GPUs.
- Budget for object-storage growth, off-site backup, and device replacement—not only the GPU.



USING AZURE VM:

Yes. **You can build this project on Microsoft Azure**, and Azure is actually a good fit for the server/cloud side.

However, I would **not move every component to Azure**. Your requirement specifically says the physical phone gateway should remain close to the phones. So I recommend a **hybrid Azure + physical phone gateway architecture**.



Requirement

Azure option

Main server

Azure VM

GPU for Qwen3

Azure GPU VM

Docker

Docker on Azure VM

PostgreSQL

Azure Database for PostgreSQL or PostgreSQL on VM

Evidence storage

Azure Blob Storage

Secrets

Azure Key Vault

Private networking

Azure VNet + VPN

Monitoring

Azure Monitor

Logs

Azure Monitor / Log Analytics

AI agent

Python application on VM/container

Local/open model

Qwen3 on GPU VM

LLM serving

vLLM

Mobile phones

Physical gateway outside/near Azure

Mobile communication

Private VPN